MANCHESTER 1824
The University of Manchester

Comparing Arm Tracking Teleoperation to Conventional Joystick Control of a Laterally Mounted Robot Arm

A dissertation submitted to The University of Manchester for the degree of

Bachelor of Engineering in Electronic Engineering

in the Faculty of Science and Engineering

Year of submission

2025

Student ID

11006231

School of Engineering

Contents

Contents 2

Abstract 4

Declaration of originality 5

Copyright statement 6

Acknowledgments 7

1 Introduction 8

1.1 Background and motivation 8

1.2 Aims and objectives 9

2 Examination of existing human to robot arm teleoperation methods 9

2.1 Why choose teleoperation over conventional input methods 9

2.2 Requirements of a suitable tracking system 10

2.3 Survey of existing tracking technologies 11

3 Comparing two viable computer vision based methods of robot arm teleoperation 12

3.1 Comparison of both computer vision methods 12

3.2 Results of joint angle tracking methods 14

3.3 Final Tracking Method 16

4 Experiments to characterise the difference in performance of arm tracking vs conventional control methods 17

4.1 Description of experiment to compare conventional vs human arm tracking 17

4.2 Scenario 1: Random position matching task 18

4.3 Scenario 2: Simple environment pick and place task 19

4.4 Scenario 3: Obstacle-rich environment pick and place task 20

5 Conclusions and future work 21

5.1 Conclusions 21

5.2 Future work 21

Appendices 22

A Project outline 22

B Risk assessment **22**

Word count: 1234

Abstract

Abstract goes here.

Declaration of originality

I hereby confirm that no portion of the work referred to in the thesis has been submitted in support of an application for another degree or qualification of this or any other university or other institute of learning.

Copyright statement

i The author of this thesis (including any appendices and/or schedules to this thesis) owns certain copyright or related rights in it (the "Copyright") and s/he has given The University of Manchester certain rights to use such Copyright, including for administrative purposes.

ii Copies of this thesis, either in full or in extracts and whether in hard or electronic copy, may be made *only* in accordance with the Copyright, Designs and Patents Act 1988 (as amended) and regulations issued under it or, where appropriate, in accordance with licensing agreements which the University has from time to time. This page must form part of any such copies made.

iii The ownership of certain Copyright, patents, designs, trademarks and other intellectual property (the "Intellectual Property") and any reproductions of copyright works in the thesis, for example graphs and tables ("Reproductions"), which may be described in this thesis, may not be owned by the author and may be owned by third parties. Such Intellectual Property and Reproductions cannot and must not be made available for use without the prior written permission of the owner(s) of the relevant Intellectual Property and/or Reproductions.

iv Further information on the conditions under which disclosure, publication and commercialisation of this thesis, the Copyright and any Intellectual Property and/or Reproductions described in it may take place is available in the University IP Policy (see `http://documents.manchester.ac.uk/DocuInfo.aspx?DocID=24420`), in any relevant Thesis restriction declarations deposited in the University Library, The University Library's regulations (see `http://www.library.manchester.ac.uk/about/regulations/`) and in The University's policy on Presentation of Theses.

Acknowledgments

Acknowledgments go here.

1 Introduction

The goal of this project is to create a system for controlling and teleoperating a robot arm with a human arm. In order to maintain simplicity in the project scope, the robot arm is assumed to be shoulder-mounted and has the seven degrees of freedom (DoF) **ref:arm_dofs**, matching that of a healthy human arm. An example of the humanoid robot arm (the YuMi robot) is shown in figure 1.

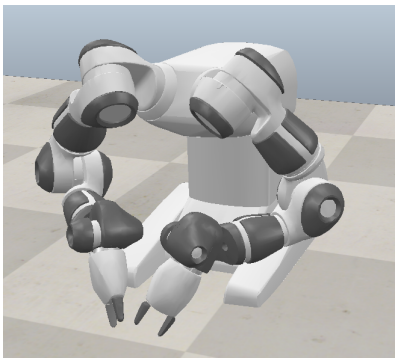

Fig. 1. YuMi Robot Arm To Be Used In Experiment

1.1 Background and motivation

The motivation for this project comes from articles within the books: **ref:Zhou_Yang_Wang_Dong_2025** **ref:Zhou_Yang_Wang_Dong_2025** and **ref:Lyu_Yang_Yang_2025** **ref:Lyu_Yang_Yang_2025**. Both of these articles use an inertial measurment unit (IMU) to gather data, but only maps the wrist position to the robot's end effector.

Redundancy in the degrees of freedom of the arm means that a valid end effector position can have multiple positions of the previous links **ref:nullspace**. Therefore this ambiguity could lead to collisions with obstacles in the environment, and lack of safety for people within the envelope of the robot arm.

Some specific use cases will require absolute control of every link of the robot arm - for example underwater sample collection. Whilst **ref:Jin_Ji_Yan_2023** describes a potential autonomous solution **ref:Jin_Ji_Yan_2023**, this method is comparatively slow to the prospect of remote teleoperation. Consequently, the complexity increase for the case where nothing is stationary, a more realistic use case, would mean that this method would be inefficient compared to the intuition of a trained operator.

Additionally, **ref:vranalysis** shows that for tasks involving teleoperation of a robot for pick and place

tasks, virtual reality style mapping of head movements with arm to end effector mapping is more efficient and less error prone than a typical joystick **ref:vranalysis**. Hence, exploration of the performance characteristics between conventional control methods and arm mapped teleoperation is undertaken.

1.2 Aims and objectives

The overall aim is to compare the performance of a system for mapping the joint angles of a human arm to a shoulder-mounted robot arm to conventional control methods.

To achieve this goal there will be an evaluation of available arm tracking methods, including:

Section 2 - Identifying and comparing existing technologies, including IMU based solutions discussed earlier, attitude and heading reference systems **ref:Mazomenos_Biswas_Cranny_Rajan_Maharatna** computer vision based systems **ref:Phuong_Cong_2024**, and flex sensor based systems **ref:Nalam_Maniva**

Section 3 - Developing a version of the tracking system suitable for controlling a simulation of the YuMi robot arm.

Section 4 - Devising experiments appropriate to comparing conventional methods of control to arm tracking teleoperation methods.

Section 5 - Analysing the difference, if any, in performance between either method of control.

2 Examination of existing human to robot arm teleoperation methods

2.1 Why choose teleoperation over conventional input methods

The measurement of human arm joint angles has been an interest of academia for several decades: with one of the first examples being demonstrated by **ref:videobasedumanmotioncapture**. In industrial settings, full body tracking has primarily been used in the film industry to accurately measure an actor's performance **ref:mocapwhy** and map them onto computer generated characters **ref:whatmocap**.

This approach could also be applied to the control of robotics: have a skilled operator remotely control industrial robots where traditional control methods (e.g. joysticks) are unsuitable. Research

has shown that this approach could lead to 1.8 times faster task completion time and with a quarter of the errors (1.54% compared to 8.65%) in the best cases **ref:vranalysis**.

Another way to increase the performance of a task with teleoperation is mapping the head movements of an operator to the camera of the robot performing the task **ref:teleopanalysis**. However, this study was based primarily on the control of wheeled robots; manual joystick controls exceeded the performance of teleoperation in this specific task.

These two cases suggests that whether teleoperation is more effective than traditional joysticks is strongly dependent on the task. Overall, in the case of robot arm control, there is potential for performance increases via human arm teleoperation **ref:Zhou_Yang_Wang_Dong_2025**, strengthening the justification for this project.

2.2 Requirements of a suitable tracking system

In order to compare the final system to existing systems and research, desired characteristics must be identified.

Accuracy requirements - Experienced surgeons have been measured to have $5°$ of vertical movement range on all joints when asked to hold their hand on a target, and close to $3°$ range for horizontal movement **ref:movementrange**. Therefore any measurement lower than this resolution can be characterised as 'noise' from human inaccuracy, and can be smoothed out.

Latency requirements - Teleoperation task performance generally degrades around the two key figures of 600 ms and 1500 ms **ref:teleoplatency**. Thus, the Nyquist criterion equivalent of this threshold is used as a well informed estimate of how much latency can be present in the system, leaving a final requirement of 300 ms.

Resource requirements - The budget for this project is £180, and 30 unit credits worth of time, equating to 240 hours of time commitment. Both of these factors will mean that the final solution favours simpler, more complete solutions - anything novel is a potential failure point that could mean the project does not complete.

2.3 Survey of existing tracking technologies

In order to fully evaluate the myriad of arm tracking methods available, a Pugh Matrix has been created to evaluate and compare several factors between each method. Each method will score $1 - 5$ in each category, with a brief comment justifying the score in each column.

Below are the decision factors of the Pugh Matrix and a brief description of how they are scored:

Complexity - This is calculated based on a well reasoned estimate of the total time to build the system. Each multiple of 48 hours reduces the score by 1, e.g. $0 - 48$ hours would score 5, $49 - 96$ hours would score 4.

Cheapness - This is calculated based on a well reasoned estimate of the total cost of building the system. Each multiple of £36 reduces the score by 1, e.g. £$0 -$ £36 would score 5, £$37 -$ £72 would score 4.

Accuracy - This is calculated from available research on the accuracy of this method. Each multiple of $5°$ above $5°$ reduces the score by 1, e.g. an absolute accuracy of $6 - 10°$ would score 4, $11 - 15°$ would score 3.

Latency - This is calculated from available research on the time it takes to obtain a single measurement of all 7 joint angles from this method. Each multiple of 100 ms above 300 ms reduces the score by 1, e.g. a latency of 301 ms $- 400$ ms would score 4.

Analysis of existing methods shows that the two computer vision approaches score similarly in terms of viability, therefore further analysis will be required to confirm the viability of each method. The accuracy advantage of the marker tracking method makes it more appealing, though this is offset by the stupendous cost of conventional Vicon camera systems.

Open source software along with ArUco markers could be used to offset this cost, which will make further analysis necessary before choosing a final method. For example, **ref:arucomarkers** describes using ArUco markers along with standard home cameras to perform specialised kinematic analysis **ref:arucomarkers**.

These will be performed in the form of accuracy and latency experiments, which will give the freedom of being able to analyse more factors than the literature - including accuracy of measuring multiple joints angles simultaneously.

Method	Complexity $1-5$	Cheapness $1-5$	Accuracy $1-5$	Latency ($1-5$)	Total
IMU-based	3 100 hours inc PCB design + drift correction	4 £$50-$ £60	3 $2.8°$ **ref:imuerror**	5 <1 ms **ref:imudatasheet**	15
Pose estimation via computer vision	5 Existing solutions reduce complexity $30-40$ hours	5 £30 USB Camera	3 $10°-15°$ **ref:cverror**	5 ≤ 100 ms **ref:cvlatency**	18
Flex Sensors	2 150 hours	4 £40-60	4 $10°$ **ref:bendaccuracy**	5 ≤ 100 ms	15
Marker tracking via computer vision	5 Very complex, though can be simplified through FOSS	3 £180+ for Vicon, cheaper solutions exist	5 <5° error	5 100-200ms	18

Table 1. Pugh Matrix with detailed scoring justifications

3 Comparing two viable computer vision based methods of robot arm teleoperation

3.1 Comparison of both computer vision methods

In order to evaluate the accuracy of both computer vision methods, a simple experiment was set up to measure the absolute accuracy on the joint angles that can be measured. A goniometer is attached to the wall at $90°$, with two end positions marked at $90°$ and $180°$. The subject will move their arm repeatedly between these markers, and the recording is captured by a camera as described in fig 2.

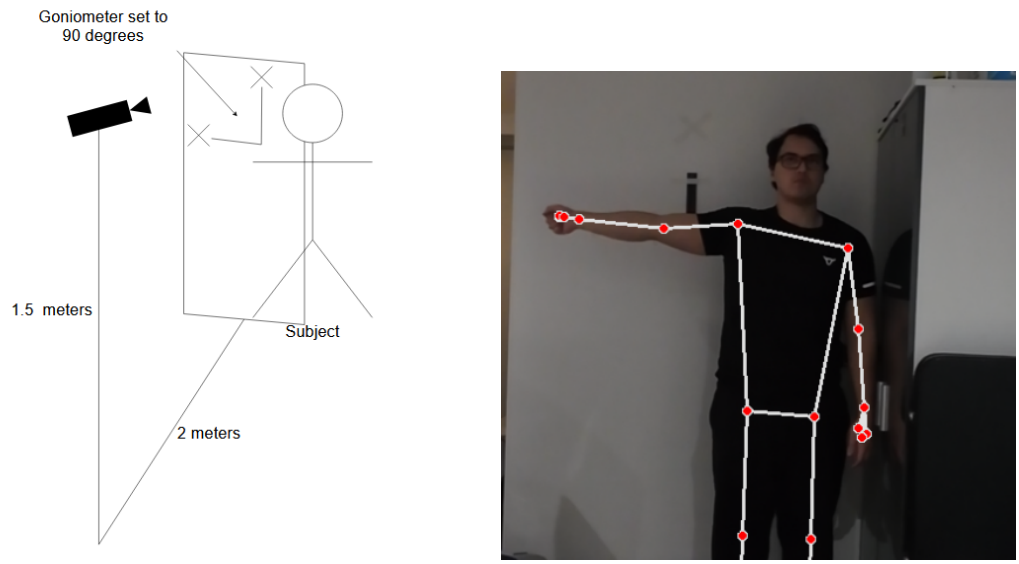

(a) A diagram of the experimental setup. **(b)** Example data capture

Fig. 2. The experimental setup, showing data capture.

The subject will then flex and extend their elbow between the two endpoints. In order to obtain a sufficient amount of data, 5 videos are recorded, each showing 10 flexion and extensions.

This creates enough samples to perform ANOVA between each video, to maintain experimental integrity. In order to measure error, the graph of elbow flexion angles is extended towards the $90°$ and $180°$ endpoints. This secondary graph is then subtracted from the first graph to obtain the error value, from which we can calculate the root mean square error, and is shown in fig 3.

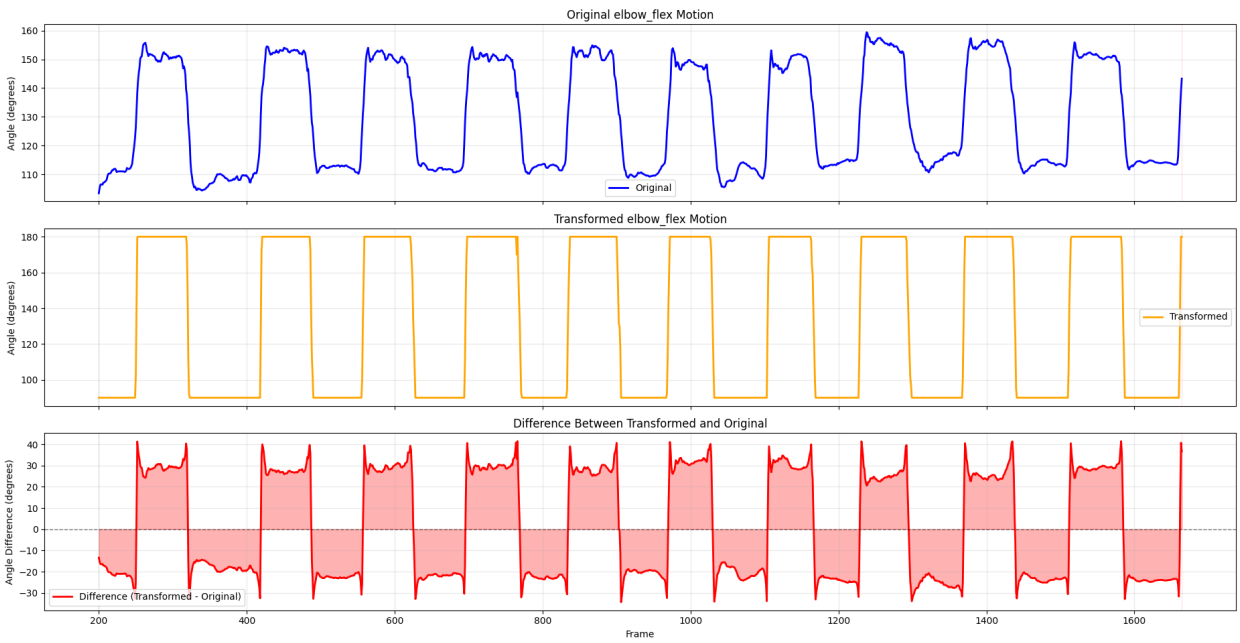

Fig. 3. Showing how RMSE is calculated from elbow flexion data.

Upon analysis of the RMSE, it was found that this value was $25.83°$. which was greater than the lit-

erature value of $10.37°$ **ref:cverror**. Frame by frame analysis of the video showed that this descrepancy was not attributable to human error - the arm was clearly flexing and extending to the correct endpoints.

Plotting the inferred joint positions frame by frame revealed the culprit - perspective! The machine learning model used for joint position extraction is finding extra movement in the depth axis of the video, and attributing it to roll in the elbow axis, instead of flexion. This is clearly shown in fig 4.

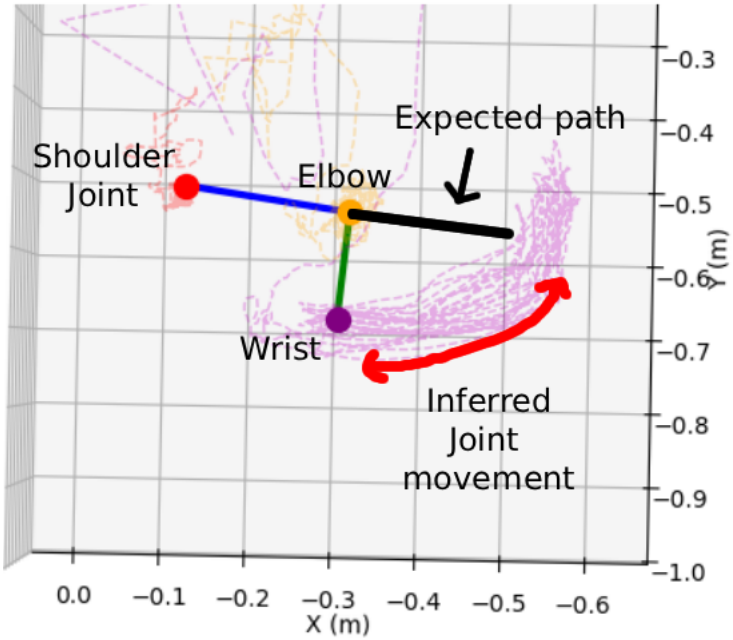

Fig. 4. Top view of inferred 3D joint positions.

To analyse the performance of joint angle tracking in the three dimensional case, a simple movement path is created. A 7 DoF manipulator is created from the DH parameters in table 2. These movements are then fed through a forward kinematics processing program to obtain a visualisation of this movement path, and this movement is recreated by the subject.

This path is shown in 5. This movement has been selected to be easily recreated, whilst showing movement in three different joint angles simultaneously. The same data processing for the two dimensional elbow flex case is used to obtain RMSE of three different joint angle movements.

3.2 Results of joint angle tracking methods

Table 3 shows the compiled RMSE errors for tracking each joint in both the two dimension elbow flexing case and the three dimension movement path.

Joint	Description	θ_i	α_i	d_i	a_i
1	Shoulder abduction/adduction	q_1	$\frac{\pi}{2}$	0	0
2	Shoulder flexion/extension	q_2	$\frac{\pi}{2}$	0	0
3	Shoulder rotation	q_3	$-\frac{\pi}{2}$	0	L_u
4	Elbow flexion/extension	q_4	$\frac{\pi}{2}$	0	0
5	Forearm rotation	q_5	$-\frac{\pi}{2}$	0	L_l
6	Wrist flexion/extension	q_6	$\frac{\pi}{2}$	0	0
7	Wrist radial/ulnar deviation	q_7	0	0	0

Table 2. Denavit-Hartenberg Parameters for 7-DoF Arm (where L_u is upper arm length and L_l is lower arm length)

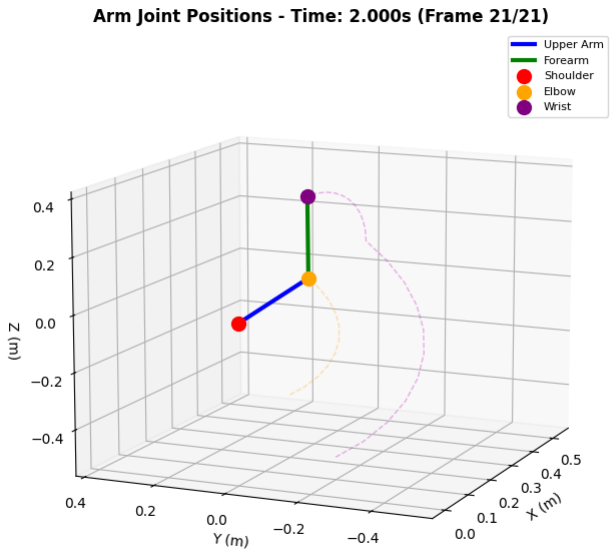

Fig. 5. 3D Joint Motion Path

Whilst the two dimensional RMSE error is larger than the literature suggests, the three dimensional joint error shows better accuracy than expected. Apart from the shoulder adduction/abduction measurements, each joint being tracked is close to the $5°$ requirements of the project.

An interesting aspect that was encountered when performing the accuracy characterisation experiments with pose estimation is shown in Fig. 6, and points to a downside not immediately apparent. The program uses machine learning to create an estimation of joint locations in an image and was ostensibly trained on a large corpus of human images. Hence, there was trouble distinguishing between real humans and pictures of humans being worn on the subject's clothing - a downside in

Method	RMSE Error Elbow Flexion 2D (°)	RMSE Error Elbow Flexion 3D (°)	RMSE Error Shoulder Flexion (°)	RMSE Error Shoulder Abduction (°)
Pose estimation	25.83	7.19	6.89	11.68
Marker Tracking				

Table 3. RMSE Error Comparison for Different Tracking Methods

the favour of the marker tracking method.

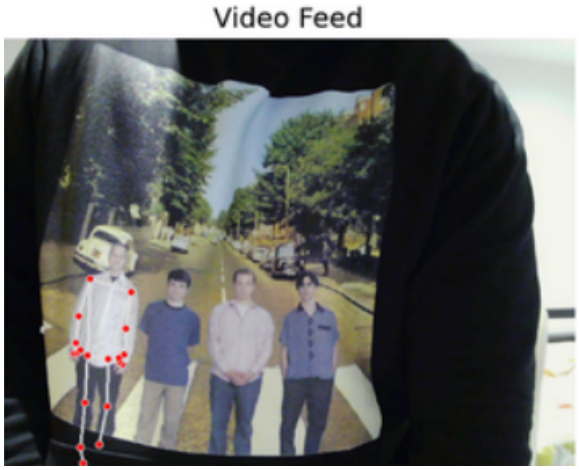

Fig. 6. Pose estimation software mistaking a picture of humans printed on the clothing for an actual person.

Despite various attempts, ArUco marker tracking seemed to be too unstable in this specific application. Both 2D paper markers and 3D printed markers were used, though the glossy surface finish of either meant that tracking was unreliable at close distances and impossible at moderate distances. This was obtained from an experiment wherein a marker is held up to the camera and rotated around various axes, measuring the captured rotation.

Fig. 7 shows a plot of the measured roll angle, and shows the gaps in which no measurement was made. Clearly this level of unreliability is unsuitable for the final tracking task, and therefore no further testing will be required.

3.3 Final Tracking Method

Results show that the pose estimation is roughly able to achieve the $5°$ measurements required for the project. Moreover, the mediapipe model used to power this is optimised for real time measurement; using a 30 fps camera, each measurement takes 33 ms, well below the required 300 ms. Hence, the comparison between conventional joystick control and human arm tracking teleoperation will use the pose estimation approach for comparison.

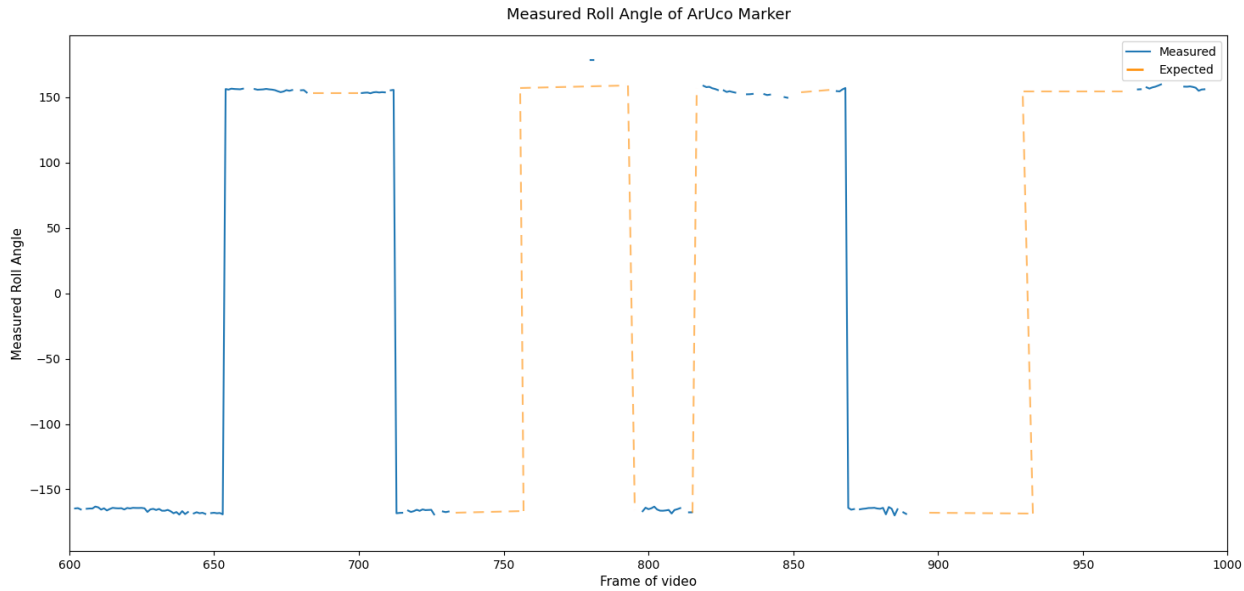

Fig. 7. Measured vs Expected graph of pitch of ArUco marker

4 Experiments to characterise the difference in performance of arm tracking vs conventional control methods

4.1 Description of experiment to compare conventional vs human arm tracking

In order to compare fairly the two methods of control, three scenarios will be examined:

Section 4.2 - Position matching task similar to **ref:movementrange ref:movementrange**. A random location within the envelope of the arm is selected and the operator must control the arm to this location as fast as possible.

Section 4.3 - Pick and place task. An object and target location spawns within the envelope of the arm, and the operator must manipulate the object to the target.

Section 4.4 - Obstacle laden pick and place task. Similar to pick and place task, but with tracking of obstacle collisions.

The operator must stand further away from the screen to be in frame of the camera for the pose estimation method. To negate the perception advantages this would incur for the controller tasks, the controller method will be performed sitting this same distance away.

Scenarios 2 and 3 consist of four different phases shown in fig 8 The object is considered 'held' as long as YuMi's grippers are within a certain distance threshold, and transition from an open state to a closed state. This prevents the experiment from being limited by the design and simulation of a suitable end effector - though this may act as an interesting extension for future research.

To measure the performance of the position matching task, the distance covered is divided by the time taken to find the normalised speed measurement. To measure the performance of the pick and place task, the cumulative time spent in each phase is measured, along with the overall time to complete the task. Likewise, to ensure fairness between each method in both pick and place tasks, seeded randomness will be used to ensure that both methods get the same trial spawn position.

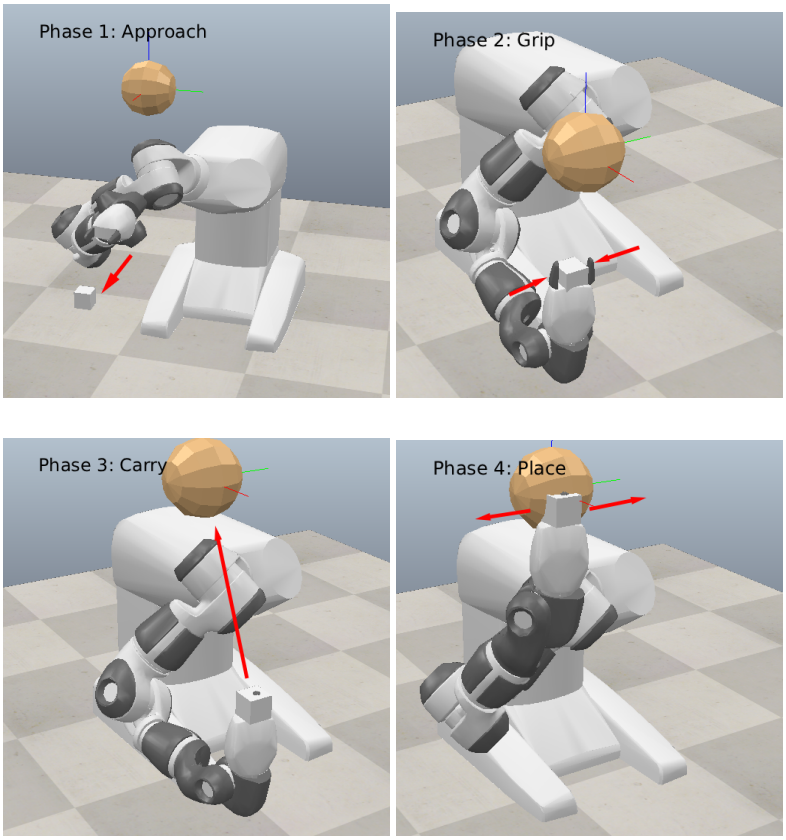

Fig. 8. Phases of the pick and place scenario

4.2 Scenario 1: Random position matching task

100 samples were taken, with 10 runs of 10 experiments. Analysis of variance (AnOVa) shows no significant deviation between each run (see appendix **TODO figure this out**). The controller method was roughly 42% faster than the pose estimation algorithm. This is likely attributable to the operator's familiarity of using controllers in video games, boasting over 10 years of meaningful experi-

ence in this field.

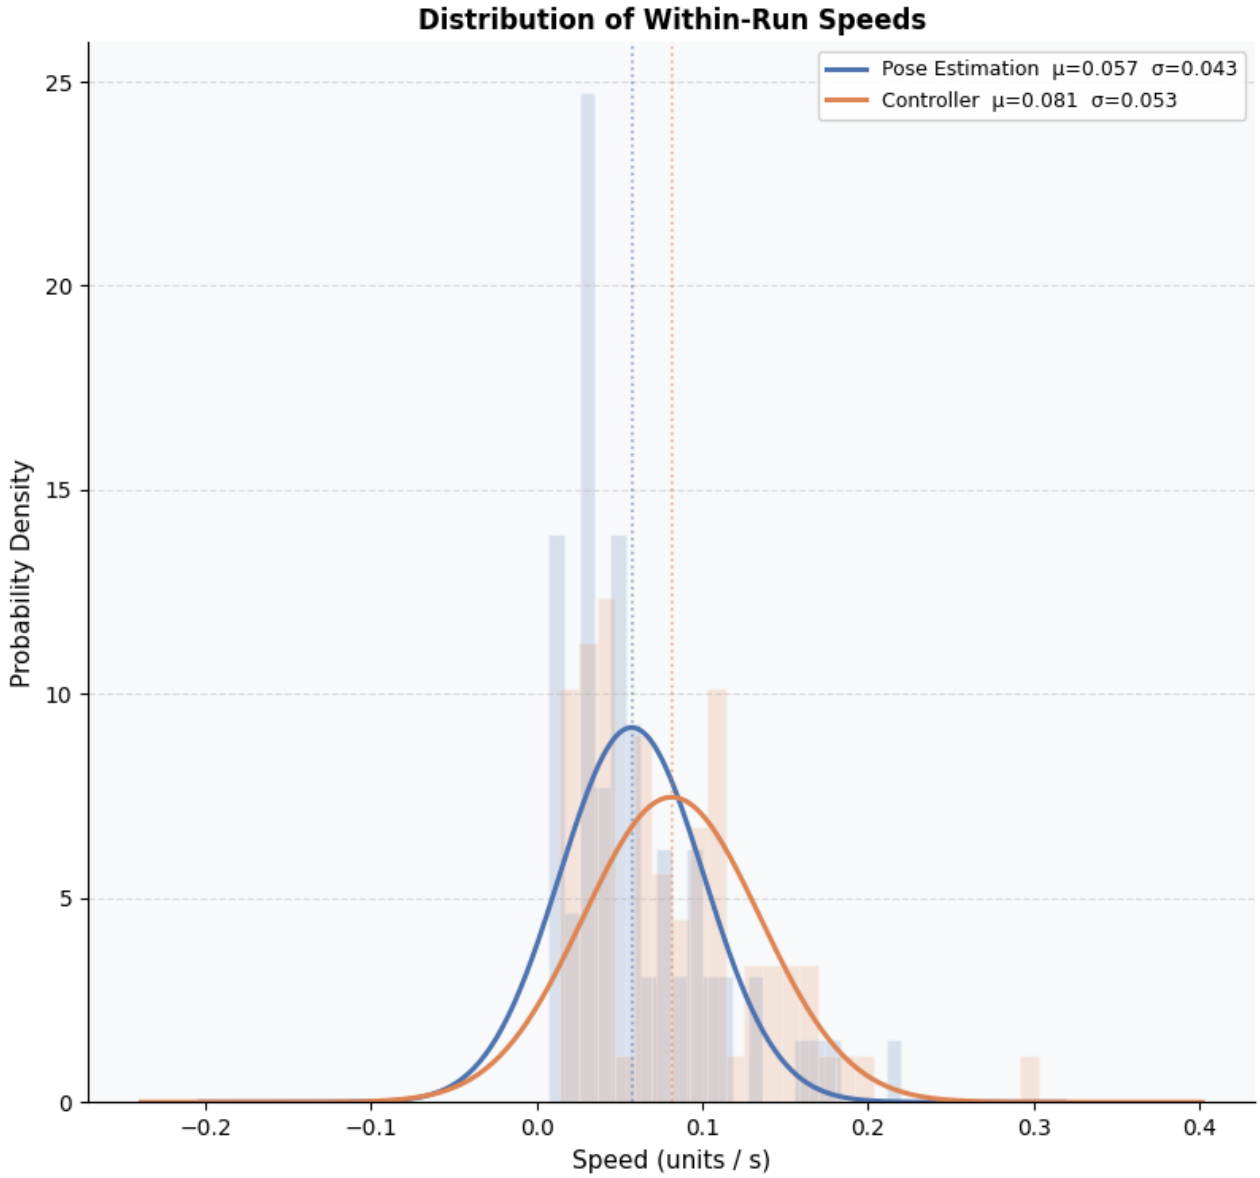

Fig. 9. Standard distribution plot of the results of either experiment

4.3 Scenario 2: Simple environment pick and place task

The pose estimation algorithm shows 60.3% reduction in time over the controller control in this task, with the noticeable difference being in the proportion of time spent in the 'approach' phase. The controller method requires the operator to orient the gripper over the object, and doing this in certain arm positions can lead to singularities in the inverse kinematics solver, hence the longer time.

These singularities can be seen in a joint angle over time plot of one of the experiments in fig 11. The advantage of the pose estimation method is that the human arm of the operator acts as a safety

net for unreachable positions.

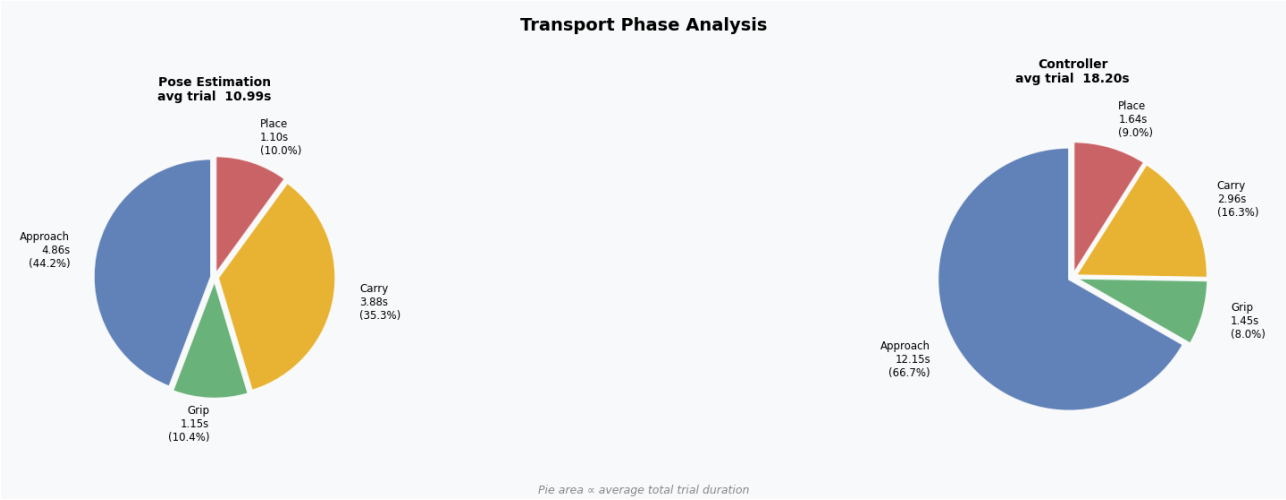

Fig. 10. Standard distribution plot of the results of either experiment

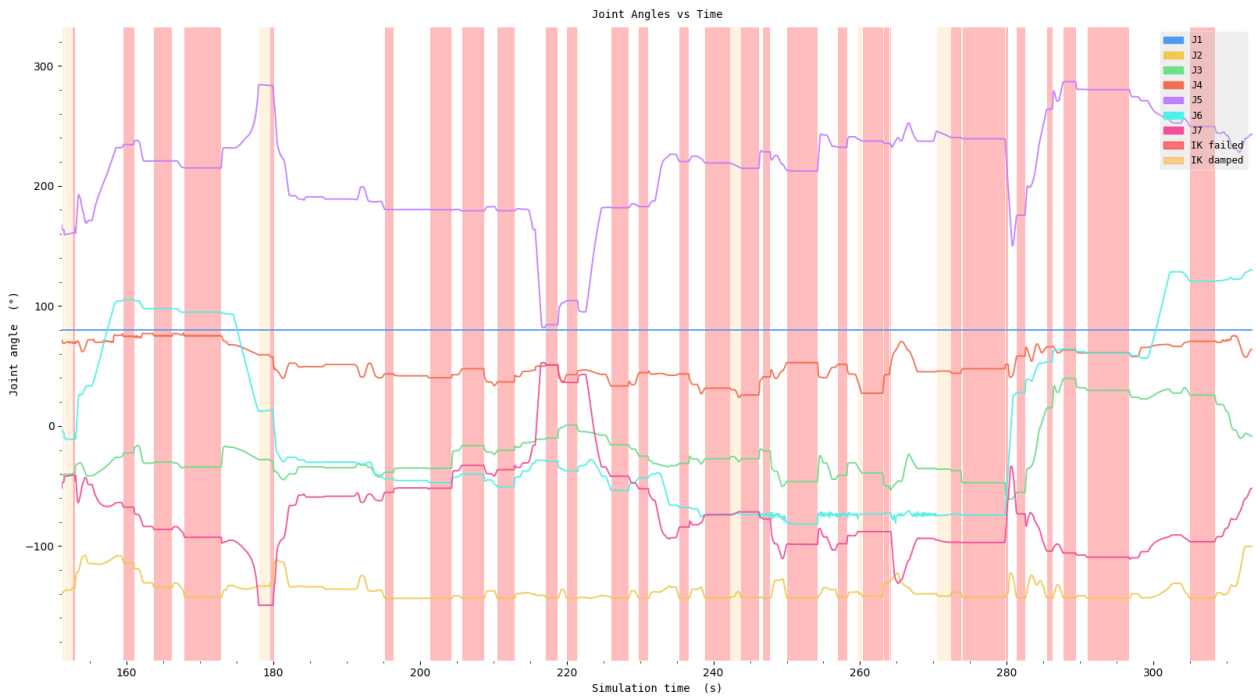

Fig. 11. Joint angle over time plot of the robot arm during a controller experiment

4.4 Scenario 3: Obstacle-rich environment pick and place task

The pose estimation algorithm is quicker on average by 57.2% in this task, likely for the same reasons as above. Additionally, the added dexterity shown led to 3.46 fewer collisions on average compared to the controller solution.

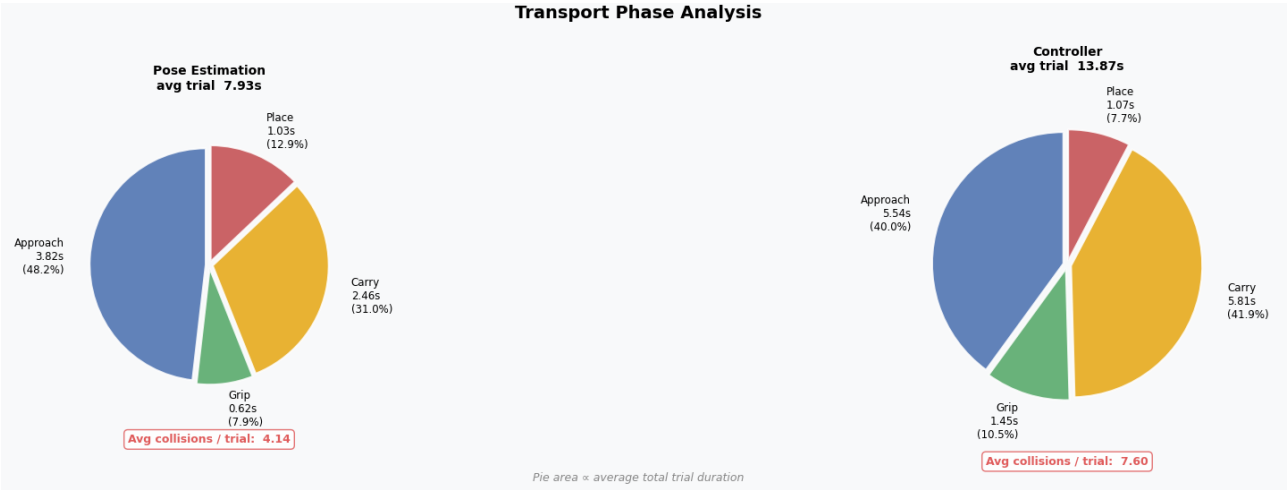

Fig. 12. Standard distribution plot of the results of either experiment

5 Conclusions and future work

5.1 Conclusions

5.2 Future work

Appendices

A Project outline

Project outline as submitted at the start of the project is a required appendix. Put here.

B Risk assessment

Risk assessment is a required appendix. Put here.